Overview

This lab does not continue the healthcare-provider scenario. Instead, you will work with data from an automobile dataset.

In this lab, you will:

- Encode ordinal categorical data
- Encode non-ordinal categorical data

About this dataset

This dataset consists of three types of entities:

1. The specification of an automobile in terms of various characteristics
2. Its assigned insurance risk rating
3. Its normalized losses in use compared to other cars

The second rating corresponds to the degree to which the automobile is riskier than its price indicates. Cars are initially assigned a risk factor symbol that's associated with its price. Then, if it's riskier (or less risky), this symbol is adjusted by moving it up (or down) the scale. Actuarians call this process *symboling*. A value of *+3* indicates that the car is risky. A value of *-3* indicates that the car is probably safe.

The third factor is the relative average loss payment per insured vehicle year. This value is normalized for all cars within a particular size classification (two-door small, station wagons, sports or speciality, and others). It represents the average loss per car per year.

Note: Several attributes in the database could be used as a *class* attribute.

Attribute information

Attribute: Attribute Range

1. symboling: -3, -2, -1, 0, 1, 2, 3.
2. normalized-losses: continuous from 65 to 256.
3. fuel-type: diesel, gas.
4. aspiration: std, turbo.

5. num-of-doors: four, two.
6. body-style: hardtop, wagon, sedan, hatchback, convertible.
7. drive-wheels: 4wd, fwd, rwd.
8. engine-location: front, rear.
9. wheel-base: continuous from 86.6 120.9.
10. length: continuous from 141.1 to 208.1.
11. width: continuous from 60.3 to 72.3.
12. height: continuous from 47.8 to 59.8.
13. curb-weight: continuous from 1488 to 4066.
14. engine-type: dohc, dohcv, l, ohc, ohcf, ohcv, rotor.
15. num-of-cylinders: eight, five, four, six, three, twelve, two.
16. engine-size: continuous from 61 to 326.
17. fuel-system: 1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi.
18. bore: continuous from 2.54 to 3.94.
19. stroke: continuous from 2.07 to 4.17.
20. compression-ratio: continuous from 7 to 23.
21. horsepower: continuous from 48 to 288.
22. peak-rpm: continuous from 4150 to 6600.
23. city-mpg: continuous from 13 to 49.
24. highway-mpg: continuous from 16 to 54.
25. price: continuous from 5118 to 45400.

Dataset attributions

This dataset was obtained from: Dua, D. and Graff, C. (2019). UCI Machine Learning Repository (http://archive.ics.uci.edu/ml). Irvine, CA: University of California, School of Information and Computer Science.

Step 1: Importing and exploring the data

You will start by examining the data in the dataset.

To get the most out of this lab, read the instructions and code before you run the cells. Take time to experiment!

Start by importing the pandas package and setting some default display options.

```
In [1]:   import pandas as pd
```

```
pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)
```

Next, load the dataset into a pandas DataFrame.

The data doesn't contain a header, so you will define those column names in a variable that's named `col_names` to the attributes listed in the dataset description.

In [2]:
```
url = "imports-85.csv"
col_names=['symboling','normalized-losses','fuel-type','aspiration','num-of-do
                    'length','width','height','curb-weight','e
                    'fuel-system','bore','stroke','compression

df_car = pd.read_csv(url,sep=',',names = col_names ,na_values="?",  header=Non
```

First, to see the number of rows (instances) and columns (features), you will use `shape`.

In [3]:
```
df_car.shape
```

Out[3]: (205, 25)

Next, examine the data by using the `head` method.

In [4]:
```
df_car.head(5)
```

Out[4]:

	symboling	normalized-losses	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	NaN	gas	std	two	convertible	rwd	front
1	3	NaN	gas	std	two	convertible	rwd	front
2	1	NaN	gas	std	two	hatchback	rwd	front
3	2	164.0	gas	std	four	sedan	fwd	front
4	2	164.0	gas	std	four	sedan	4wd	front

There are 25 columns. Some of the columns have numerical values, but many of them contain text.

To display information about the columns, use the `info` method.

```
In [5]:  df_car.info()

         <class 'pandas.core.frame.DataFrame'>
         RangeIndex: 205 entries, 0 to 204
         Data columns (total 25 columns):
          #   Column             Non-Null Count  Dtype
         ---  ------             --------------  -----
          0   symboling          205 non-null    int64
          1   normalized-losses  164 non-null    float64
          2   fuel-type          205 non-null    object
          3   aspiration         205 non-null    object
          4   num-of-doors       203 non-null    object
          5   body-style         205 non-null    object
          6   drive-wheels       205 non-null    object
          7   engine-location    205 non-null    object
          8   wheel-base         205 non-null    float64
          9   length             205 non-null    float64
          10  width              205 non-null    float64
          11  height             205 non-null    float64
          12  curb-weight        205 non-null    int64
          13  engine-type        205 non-null    object
          14  num-of-cylinders   205 non-null    object
          15  engine-size        205 non-null    int64
          16  fuel-system        205 non-null    object
          17  bore               201 non-null    float64
          18  stroke             201 non-null    float64
          19  compression-ratio  205 non-null    float64
          20  horsepower         203 non-null    float64
          21  peak-rpm           203 non-null    float64
          22  city-mpg           205 non-null    int64
          23  highway-mpg        205 non-null    int64
          24  price              201 non-null    float64
         dtypes: float64(11), int64(5), object(9)
         memory usage: 40.2+ KB
```

To make it easier to view the dataset when you start encoding, drop the columns
that you won't use.

```
In [6]:  df_car.columns
```

```
Out[6]:  Index(['symboling', 'normalized-losses', 'fuel-type', 'aspiration', 'num-of-d
         oors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'lengt
         h', 'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'eng
         ine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepowe
         r', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price'], dtype='object')
```

```
In [7]:  df_car = df_car[[ 'aspiration', 'num-of-doors',  'drive-wheels',  'num-of-cyli
```

You now have four columns. These columns all contain text values.

```
In [8]:  df_car.head()
```

	aspiration	num-of-doors	drive-wheels	num-of-cylinders
0	std	two	rwd	four
1	std	two	rwd	four
2	std	two	rwd	six
3	std	four	fwd	four
4	std	four	4wd	five

Most machine learning algorithms require inputs that are numerical values.

- The **num-of-cylinders** and **num-of-doors** features have an ordinal value. You could convert the values of these features into their numerical counterparts.
- However, **aspiration** and **drive-wheels** don't have an ordinal value. These features must be converted differently.

You will explore the ordinal features first.

Step 2: Encoding ordinal features

In this step, you will use a mapper function to convert the ordinal features into ordered numerical values.

Start by getting the new column types from the DataFrame:

```
In [9]: df_car.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 4 columns):
 #   Column            Non-Null Count  Dtype
---  ------            --------------  -----
 0   aspiration        205 non-null    object
 1   num-of-doors      203 non-null    object
 2   drive-wheels      205 non-null    object
 3   num-of-cylinders  205 non-null    object
dtypes: object(4)
memory usage: 6.5+ KB
```

First, determine what values the ordinal columns contain.

Starting with the **num-of-doors** feature, you can use `value_counts` to discover the values.

```
In [10]: df_car['num-of-doors'].value_counts()
```

```
Out[10]: num-of-doors
         four     114
         two       89
         Name: count, dtype: int64
```

This feature only has two values: *four* and *two*. You can create a simple mapper that contains a dictionary:

```
In [11]: door_mapper = {"two": 2,
                        "four": 4}
```

You can then use the `replace` method from pandas to generate a new numerical column based on the **num-of-doors** column.

```
In [25]: pd.set_option('future.no_silent_downcasting', True)
         df_car['doors'] = df_car["num-of-doors"].replace(door_mapper)
```

When you display the DataFrame, you should see the new column on the right. It contains a numerical representation of the number of doors.

```
In [13]: df_car.head()
```

Out[13]:

	aspiration	num-of-doors	drive-wheels	num-of-cylinders	doors
0	std	two	rwd	four	2.0
1	std	two	rwd	four	2.0
2	std	two	rwd	six	2.0
3	std	four	fwd	four	4.0
4	std	four	4wd	five	4.0

Repeat the process with the **num-of-cylinders** column.

First, get the values.

```
In [14]: df_car['num-of-cylinders'].value_counts()
```

```
Out[14]: num-of-cylinders
         four      159
         six        24
         five       11
         eight       5
         two         4
         three       1
         twelve      1
         Name: count, dtype: int64
```

Next, create the mapper.

```
In [15]:  cylinder_mapper = {"two":2,
                             "three":3,
                             "four":4,
                             "five":5,
                             "six":6,
                             "eight":8,
                             "twelve":12}
```

Apply the mapper by using the `replace` method.

```
In [24]:  pd.set_option('future.no_silent_downcasting', True)
          df_car['cylinders'] = df_car['num-of-cylinders'].replace(cylinder_mapper)
```

```
In [17]:  df_car.head()
```

Out[17]:

	aspiration	num-of-doors	drive-wheels	num-of-cylinders	doors	cylinders
0	std	two	rwd	four	2.0	4
1	std	two	rwd	four	2.0	4
2	std	two	rwd	six	2.0	6
3	std	four	fwd	four	4.0	4
4	std	four	4wd	five	4.0	5

For more information about the `replace` method, see pandas.DataFrame.replace in the pandas documentation.

Step 3: Encoding non-ordinal categorical data

In this step, you will encode non-ordinal data by using the `get_dummies` method from pandas.

The two remaining features are not ordinal.

According to the attribute description, the following values are possible:

- aspiration: std, turbo.
- drive-wheels: 4wd, fwd, rwd.

You might think that the correct strategy is to convert these values into numerical values. For example, consider the **drive-wheels** feature. You could use *4wd = 1*, *fwd = 2*, and *rwd = 3*. However, *fwd* isn't less than *rwd*. These values don't have an

order, but you just introduced an order to them by assigning these numerical values.

The correct strategy is to convert these values into *binary features* for each value in the original feature. This process is often called *one-hot encoding* in machine learning, or *dummying* in statistics.

pandas provides a `get_dummies` method, which converts the data into binary features. For more information, see pandas.get_dummies in the pandas documentation.

According to the attribute description, **drive-wheels** has three possible values.

```
In [18]:  df_car['drive-wheels'].value_counts()
```

```
Out[18]:  drive-wheels
          fwd      120
          rwd       76
          4wd        9
          Name: count, dtype: int64
```

Use the `get_dummies` method to add new binary features to the DataFrame.

```
In [19]:  df_car = pd.get_dummies(df_car,columns=['drive-wheels'])
```

```
In [20]:  df_car.head()
```

Out[20]:

	aspiration	num-of-doors	num-of-cylinders	doors	cylinders	drive-wheels_4wd	drive-wheels_fwd	whee
0	std	two	four	2.0	4	False	False	
1	std	two	four	2.0	4	False	False	
2	std	two	six	2.0	6	False	False	
3	std	four	four	4.0	4	False	True	
4	std	four	five	4.0	5	True	False	

When you examine the dataset, you should see three new columns on the right:

- **drive-wheels_4wd**
- **drive-wheels_fwd**
- **drive-wheels_rwd**

The encoding was straightforward. If the value in the **drive-wheels** column is *4wd*, then a *1* is the value in the **drive-wheels_4wd** column. A *0* is the value for the

other columns that were generated. If the value in the **drive-wheels** column is *fwd*, then a *1* is the value in the **drive-wheels_fwd** column, and so on.

These binary features enable you to express the information in a numerical way, without implying any order.

Examine the final column that you will encode.

The data in the **aspiration** column only has two values: *std* and *turbo*. You could encode this column into two binary features. However, you could also ignore the *std* value and record whether it's *turbo* or not. To do this, you would still use the `get_dummies` method, but specify `drop_first` as *True*.

```
In [21]:  df_car['aspiration'].value_counts()
```

```
Out[21]:  aspiration
          std       168
          turbo      37
          Name: count, dtype: int64
```

```
In [22]:  df_car = pd.get_dummies(df_car,columns=['aspiration'], drop_first=True)
```

```
In [23]:  df_car.head()
```

Out[23]:

	num-of-doors	num-of-cylinders	doors	cylinders	drive-wheels_4wd	drive-wheels_fwd	drive-wheels_rwd	asp
0	two	four	2.0	4	False	False	True	
1	two	four	2.0	4	False	False	True	
2	two	six	2.0	6	False	False	True	
3	four	four	4.0	4	False	True	False	
4	four	five	4.0	5	True	False	False	

Challenge task: Go back to the beginning of this lab, and add other columns to the dataset. How would you encode the values of each column? Update the code to include some of the other features.

Congratulations!

You have completed this lab, and you can now end the lab by following the lab guide instructions.

E0123047